

ex2 - part 1

Eitan kerzhner 205697139
itamar Shachen tov 207497769

Model's Params:

POS:

- o Dropout: $p=0.5$ (default), applied after Tanh activation.
- o Optimizer: Adam.
- o Learning rate: $1e-3$ (Adam's default).
- o Hidden layer size: 128 units.
- o Word embedding size: 50 dimensions.
- o Context window: 5 words.
- o Batch size (train & dev): 64.
- o Number of epochs: 12.

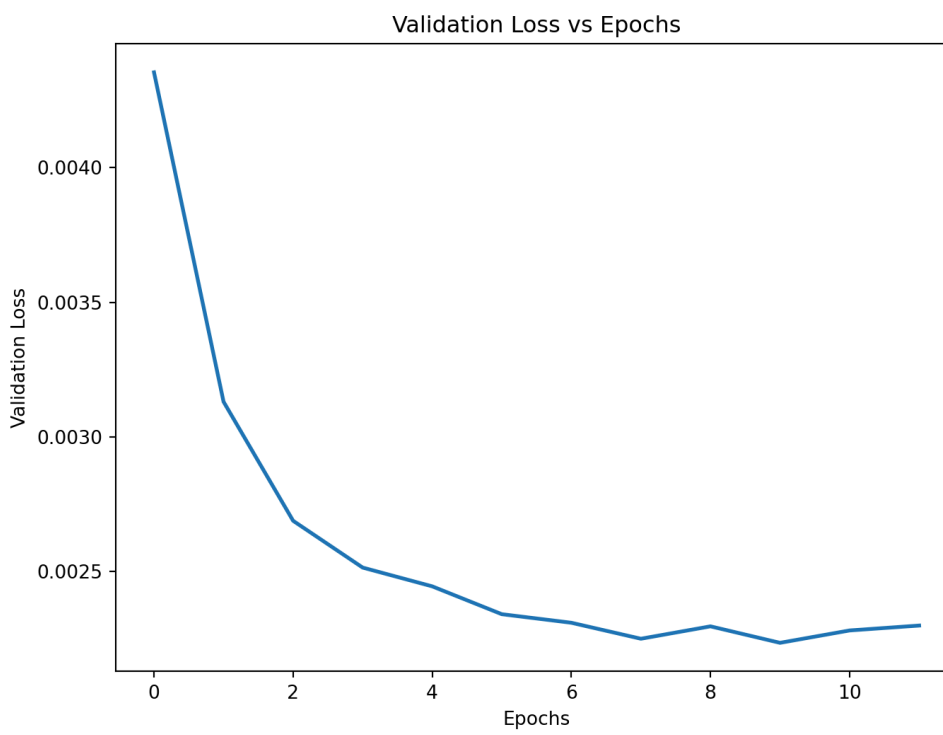
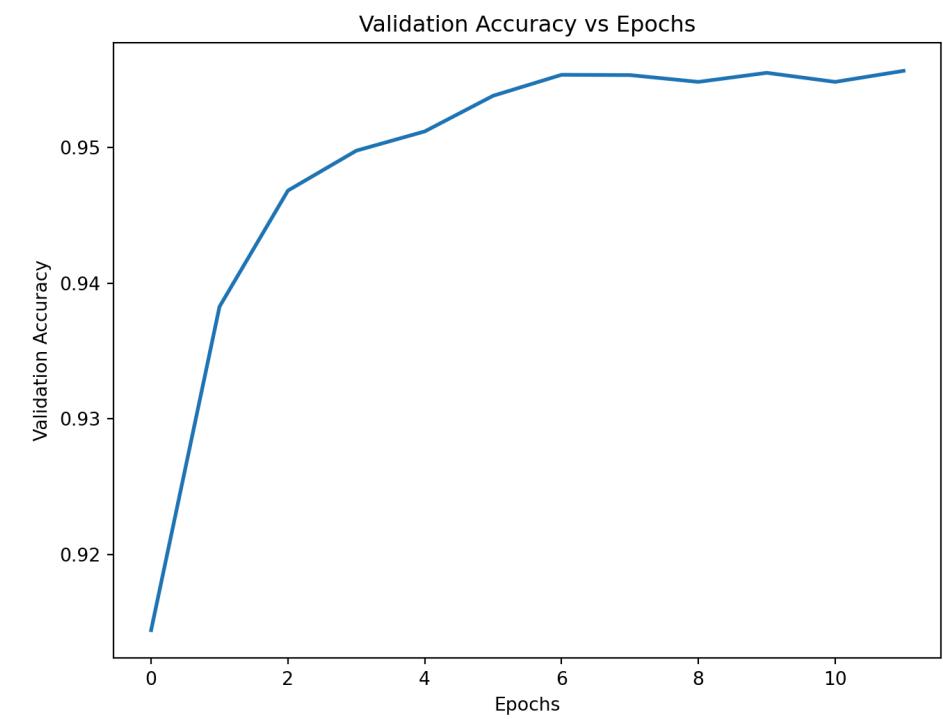
NER:

- o Dropout: $p=0.5$ (default), applied after Tanh activation.
- o Optimizer: Adam.
- o Learning rate: $1e-3$ (Adam's default).
- o Hidden layer size: 128 units.
- o Word embedding size: 50 dimensions.
- o Context window: 5 words.
- o Batch size (train & dev): 64.
- o Number of epochs: 12.

Considerations:

- o Unseen/rare words (frequency ≤ 1 in train) are mapped to an "<unknown>" token. This token has its own embedding, trained by treating these rare training words as "<unknown>".
- o To handle context for initial words, two "<s>" (start-of-sentence) tokens are prepended to each sentence. These tokens have learnable embeddings and provide left context.
- o Similarly, two "<e>" (end-of-sentence) tokens are appended for the right context of final words, also with their own learnable embeddings.

POS:



NER:

